

Session 4: CNN Hands-on Lab

Building and Training CNNs for EO Image Classification

Stylianos Kotsopoulos
EU-Philippines CoPhil Programme

Session Overview

Duration: 2.5 hours

Type: Intensive hands-on lab

Goal: Turn CNN theory into a working model

You will: - Prepare the EuroSAT dataset - Build a CNN from scratch in TensorFlow/Keras - Train with GPU acceleration in Colab - Evaluate with confusion matrix and F1 - Apply transfer learning (ResNet50)

Prerequisites: - Session 3 complete (CNN basics) - Colab account with GPU enabled - Python & NumPy fundamentals

Notebook:

[session4_cnn_classification_STUDENT.ipynb](#)

Setup & Data Preparation

Colab GPU + Environment

- Steps:** 1. Runtime → Change runtime type → Hardware accelerator → GPU
2. Verify with `!nvidia-smi`
3. `pip install tensorflow` (if needed)
4. Set seeds for reproducibility

```
1 import tensorflow as tf, numpy as np, random, os
2 SEED = 42
3 random.seed(SEED); np.random.seed(SEED); tf.random.set_seed(SEED)
4 print(tf.__version__, tf.config.list_physical_devices('GPU'))
```

EuroSAT Dataset (Sentinel-2, 10 classes)

- ~27k RGB chips (64×64) derived from S2
- Classes: AnnualCrop, Forest, Herbaceous, Highway, Industrial, Pasture, PermanentCrop, Residential, River, SeaLake

```
1 # Example TFDS approach
2 import tensorflow_datasets as tfds
3 (ds_train, ds_val, ds_test), meta = tfds.load(
4     'euosat/rgb', split=['train[:70%]', 'train[70%:85%]', 'train[85%:]'],
5     as_supervised=True, with_info=True)
```

Preprocessing: Normalize to [0,1], one-hot labels, split 70/15/15

Building a CNN from Scratch

Architecture (reference)

```
1 Input (64×64×3)
2 → [Conv(32, 3×3) + ReLU] → MaxPool
3 → [Conv(64, 3×3) + ReLU] → MaxPool
4 → [Conv(128, 3×3) + ReLU] → MaxPool
5 → Flatten → Dropout(0.5)
6 → Dense(128) + ReLU → Dropout(0.5)
7 → Dense(10) + Softmax
```

```
1 from tensorflow.keras import Sequential
2 from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Dropout, Flatten
3
4 model = Sequential([
5     Conv2D(32, (3,3), activation='relu', input_shape=(64,64,3)),
6     MaxPooling2D(),
7     Conv2D(64, (3,3), activation='relu'),
8     MaxPooling2D(),
9     Conv2D(128, (3,3), activation='relu'),
10    MaxPooling2D(),
11    Flatten(), Dropout(0.5),
12    Dense(128, activation='relu'), Dropout(0.5),
13    Dense(10, activation='softmax')
14 ])
```

Compile & Train

```
1 model.compile(optimizer='adam',
2               loss='sparse_categorical_crossentropy',
3               metrics=['accuracy', 'top_k_categorical_accuracy'])
4
5 cb = [
6     tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True),
7     tf.keras.callbacks.ModelCheckpoint('best_model.h5', save_best_only=True),
8     tf.keras.callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3)
9 ]
10
11 history = model.fit(ds_train.batch(64).prefetch(2),
12                    validation_data=ds_val.batch(64).prefetch(2),
13                    epochs=30, callbacks=cb)
```

Healthy curves: train↓, val↓ then plateau; small gap

Overfitting: large gap → add dropout/augmentation

Evaluation & Error Analysis

Accuracy + Confusion Matrix

```
1 import numpy as np, matplotlib.pyplot as plt
2 from sklearn.metrics import confusion_matrix, classification_report
3
4 y_true, y_pred = [], []
5 for x, y in ds_test.batch(64):
6     p = model.predict(x, verbose=0).argmax(axis=1)
7     y_pred.extend(p); y_true.extend(y.numpy())
8
9 cm = confusion_matrix(y_true, y_pred)
10 print(classification_report(y_true, y_pred))
```

Look for: - Which pairs are confused? (e.g., AnnualCrop vs Herbaceous)
- Per-class precision/recall balance

Visualize Misclassifications

1 # Show a grid of wrong predictions for qualitative review

- Investigate systematic errors
- Adjust augmentation / architecture accordingly

Data Augmentation

EO-aware augmentations

```
1 import tensorflow as tf
2 from tensorflow.keras import layers
3
4 data_aug = tf.keras.Sequential([
5     layers.RandomRotation(0.25),
6     layers.RandomFlip('horizontal_and_vertical'),
7     layers.RandomZoom(0.1),
8     layers.RandomContrast(0.1)
9 ])
```

- Rotations/flips OK for overhead imagery
- Brightness/contrast for atmospheric
- Avoid orientation-critical tasks if sensitive (roads)

Transfer Learning (ResNet50)

Feature extraction → fine-tuning

```
1 from tensorflow.keras.applications import ResNet50
2 from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, Dropout, Input
3 from tensorflow.keras.models import Model
4
5 base = ResNet50(include_top=False, weights='imagenet', input_shape=(64,64,3))
6 for l in base.layers[:int(0.8*len(base.layers))]:
7     l.trainable = False
8
9 x = GlobalAveragePooling2D()(base.output)
10 x = Dense(256, activation='relu')(x)
11 x = Dropout(0.5)(x)
12 out = Dense(10, activation='softmax')(x)
13 model = Model(base.input, out)
```

Strategy: - Start with frozen base → quick convergence
- Unfreeze top blocks for small accuracy gains

Compare Approaches

Approach	Train Time	Accuracy
From scratch	15–25 min	92–95%
Feature extraction	5–10 min	94–96%
Partial fine-tune	10–20 min	95–97%

Tip: Use the fastest path during live sessions, fine-tune offline later

Troubleshooting

Common issues & fixes

- **GPU not detected:** set runtime → GPU; restart runtime
- **OOM error:** lower batch size; fewer filters; mixed precision
- **Accuracy stuck ~10%:** check labels; learning rate; normalization
- **Overfitting:** stronger augmentation; more dropout; L2; early stop
- **Slow training:** reduce model depth; use caching/prefetch

Philippine Context

Why CNNs matter operationally

- National land cover refresh (PhilSA)
- Cloud masking for S2 mosaics
- Disaster mapping (flood/damage)
- Urban growth monitoring
- Supports DENR, DA, NDRRMC, LGUs

Next steps: move to Day 3 (U-Net segmentation, flood mapping)

Summary & Notebook

What you achieved today

1. Built and trained a CNN classifier (EuroSAT)
2. Evaluated with robust metrics and confusion matrix
3. Applied transfer learning for higher accuracy
4. Learned practical debugging strategies

Notebook:

`session4_cnn_classification_STUDENT.ipynb`

Render slides (local):

```
1 cd course_site/day2/presentations
2 quarto render session4_cnn_lab.qmd
```